

Beyond Software Defenses: Hardware-assisted CFI on Arm-based systems

Michalis Pappas
mpappas@fastmail.fm

HKOSCon 2024

About

Michalis Pappas

Staff Engineer at Unikraft GmbH

<https://github.com/michpappas>



Motivation: Hardware-assisted CFI on Unikraft for arm64



michpappas commented on Dec 12, 2021 • edited ▾

Member ...

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☐ Updated relevant documentation.

Base target

- Architecture(s): arm64
- Platform(s): common
- Application(s): N/A

Additional configuration

This PR introduces `CONFIG_ARM64_FEAT_PAUTH` to enable Pointer Authentication support.

Description of changes

Pointer Authentication (PAuth) allows signing and authenticating pointers to harden the system against classes of attacks that rely on the manipulation of pointers, such as ROP. Pointer Authentication Codes (PACs) are created by signing a pointer and a 64-bit modifier with an 128-bit key. The modifier is a value that is normally used to restrict the PAC to a specific context (eg the SP). Keys are stored in registers.

The PAC is stored in the unused upper bits of the memory address, and can be later verified to ensure that the pointer has not been tampered.

The reference algorithm is [the QARMA block cipher](#), but it is possible for architectures to use an IMPLEMENTATION DEFINED algorithm instead.



michpappas commented on Mar 6, 2022

Member ...

Prerequisite checklist

- ☒ Read the [contribution guidelines](#) regarding submitting new changes to the project;
- ☒ Tested your changes against relevant architectures and platforms;
- ☒ Ran the [checkpatch.pl](#) on your commit series before opening this PR;
- ☐ Updated relevant documentation.

Base target

- Architecture(s): [e.g. `arm64`]
- Platform(s): [common]
- Application(s): [N/A]

Additional configuration

This PR introduces a new Kconfig option, `CONFIG_ARM64_FEAT_BTI`. For details, see below.

Description of changes

`FEAT_BTI` is a hardware protection against JOP-like attacks.

To do that, it introduces:



- The BTI instruction.
- The GP field in Stage 1 PTES.
- The `PSTATE.BTYPE` field.




BTI instructions, aka landing pads, are placed by the compiler at branch targets. On runtime, branches that do not land on a BTI instruction trigger an Branch Target Exception.


The GP field indicates whether a page is guarded with BTI. This is allows backwards compatibility, by disabling BTI on pages that contain non-BTI guarded code. Notice that BTI instructions on unguarded pages execute as NOP.





`PSTATE.BTYPE` encodes the type of an indirect jump, ie the branch instruction, the registers used to carry parameters, and whether the target page is guarded or not. When an indirect branch is taken, the processor checks whether `PSTATE.BTYPE` matches the type of the branch target, and on negative match it generates an Branch Target Exception. The purpose of this is to further limit the scope of possible gadgets among BTI protected branches. Notice that there are exceptions to this. For details see D5.4.4 in [2].





(Non-exhaustive) evolution₊ of CFI attacks & mitigations




 **1988:** The Morris worm exploits a buffer overflow in the wild
 1996: "Smashing The Stack For Fun And Profit" Aleph One, Phrack 49




 **1997:** Solaris switches to non-executable stacks
 1997: "Getting around non-executable stack (and fix)" Solar Designer
BUGTRAQ
 1998: FreeBSD switches to non-executable stacks



 **1998:** "StackGuard: Automatic Adaptive Detection and Prevention of Buffer-Overflow Attacks" Conowan et al, USENIX




 **1999:** Stack Shield: A "stack smashing" technique protection tool for Linux
 2000: "Bypassing StackGuard and StackShield" Bulba and Kil3r, Phrack 52
 2003: "ProPolice: GCC Extension for Protecting Applications from Stack-Smashing Attacks" Hiroaki Etoh IBM
 2004: GCC 3.x introduces stack protector (SSP) based on ProPolice

 **2001:** PaX introduces ASRL
 2002: "Bypassing PaX ASLR protection" Tyler Durden, Phrack 59
 2003: OpenBSD adds ASLR
 2005: Linux adds ASLR based on PaX

 **2003:** AMD introduces NX bit
 2004: Intel adds NX bit
 2007: Arm introduces WXN

 **2008:** Return-oriented programming (RoP) / Jump-oriented programming (JoP)
 "The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86 Architecture)" Hovav Shacham
 "Return-Oriented Programming: Exploits Without Code Injection" Stefan Savage et al

 **2013:** PaX introduces KASLR
 Attackers have demonstrated numerous ways to bypass KASLR (<https://github.com/bcoles/kasld>)

 **2016:** Arm and Intel announce hardware-assisted CFI
 Arm: PAuth and BTI
 Intel: CET & IBT

Hardware-assisted Control Flow Integrity (CFI)

Hardware-assisted CFI in Arm:

- Pointer Authentication (PAuth): Enforces CFI by authenticating pointers.
- Branch Target Identification (BTI): Enforces CFI by restricting indirect jumps.

Threat model:

- Attacker has arbitrary READ & WRITE and is aiming for arbitrary EXEC.
- Non-executable stack in place.

Pointer Authentication (PAuth)

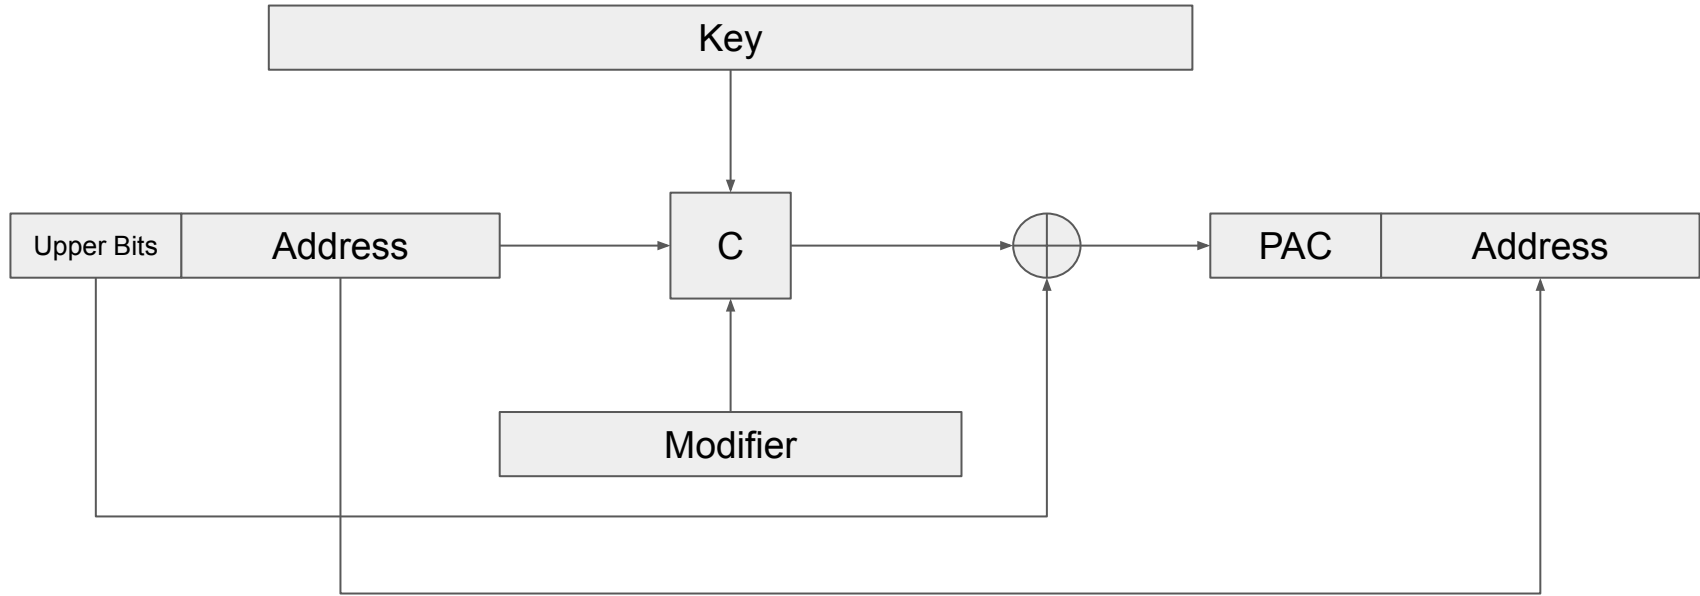
Pointer Authentication (PAuth)

Based on cryptographic signing and authentication of pointers.

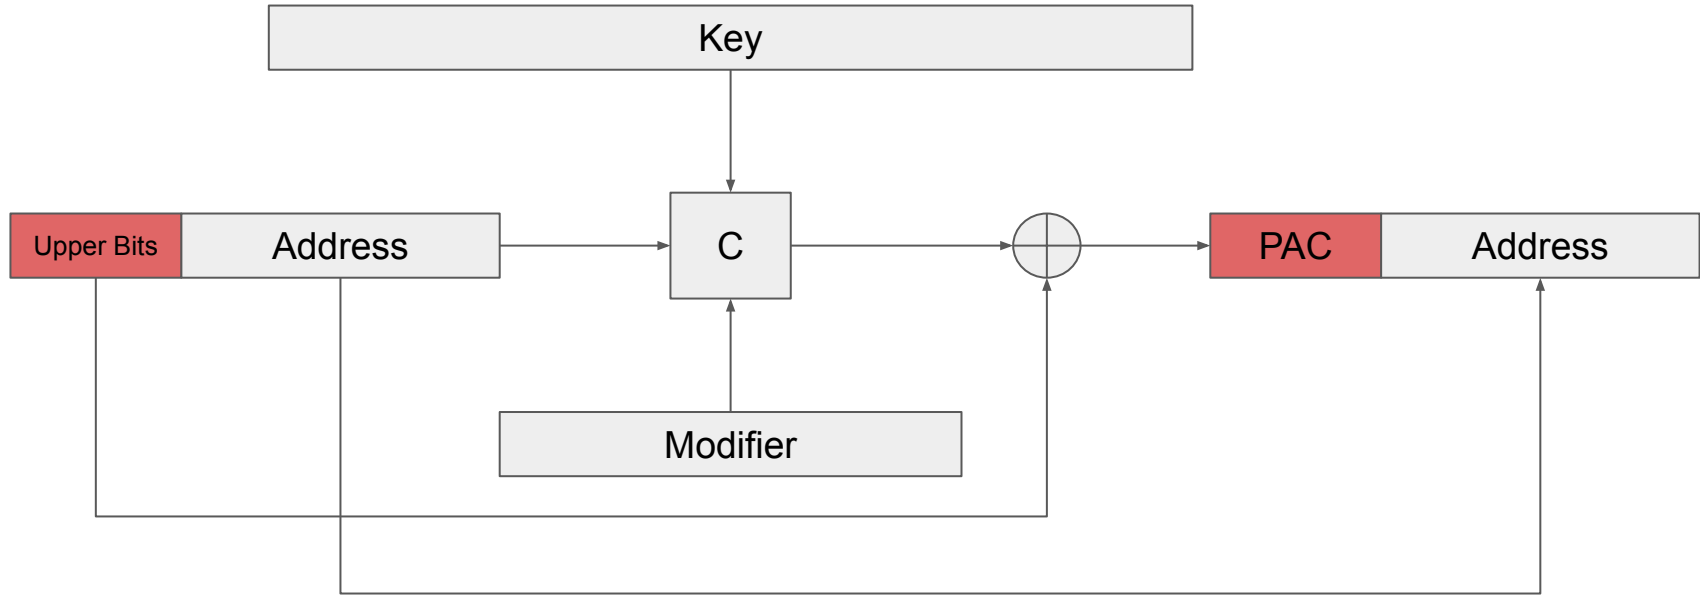
- **Armv8.3-A** introduces **FEAT_PAuth** as a **MANDATORY** extension.
- **Armv8.6-A** introduces **FEAT_PAuth2** as a **MANDATORY** extension.

In this talk we will be focusing on **FEAT_PAuth2**.





PAuth Operation: Sign







PAAuth: Upper Bits / PAC



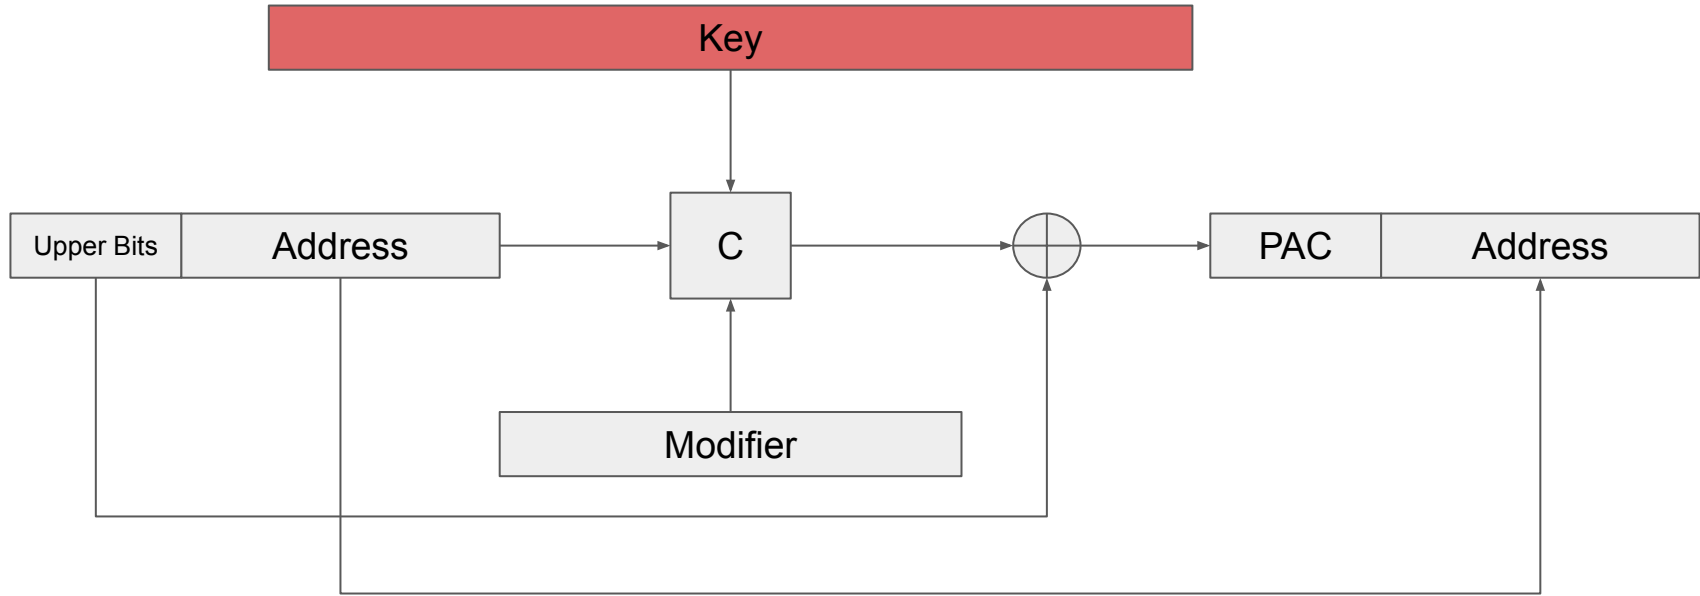
PAuth: Upper Bits / PAC



PAuth: Upper Bits / PAC



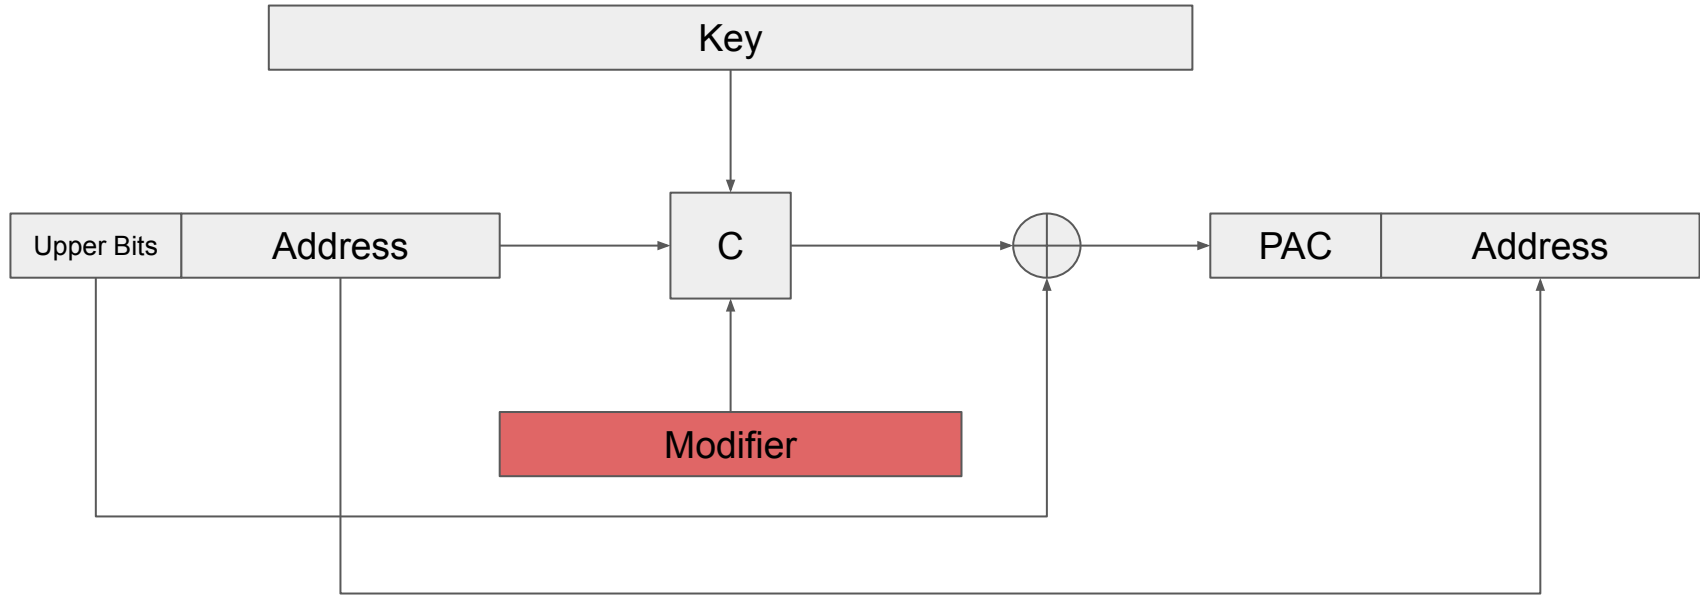
PAuth: Key



PAuth: Key Registers

Key	Purpose	Mnemonic
APIAKey_ELx	Instruction signing Key A	IA
APIBKey_ELx	Instruction signing Key B	IB
APDAKey_ELx	Data signing Key A	DA
APDBKey_ELx	Data signing Key B	DB
APGAKey_ELx	Generic signing key	G

PAuth: Modifier



PAuth: Modifier

Problem:

- Pointer substitution attacks.
- An attacker can reuse signed pointers (remember threat model).

Mitigation:

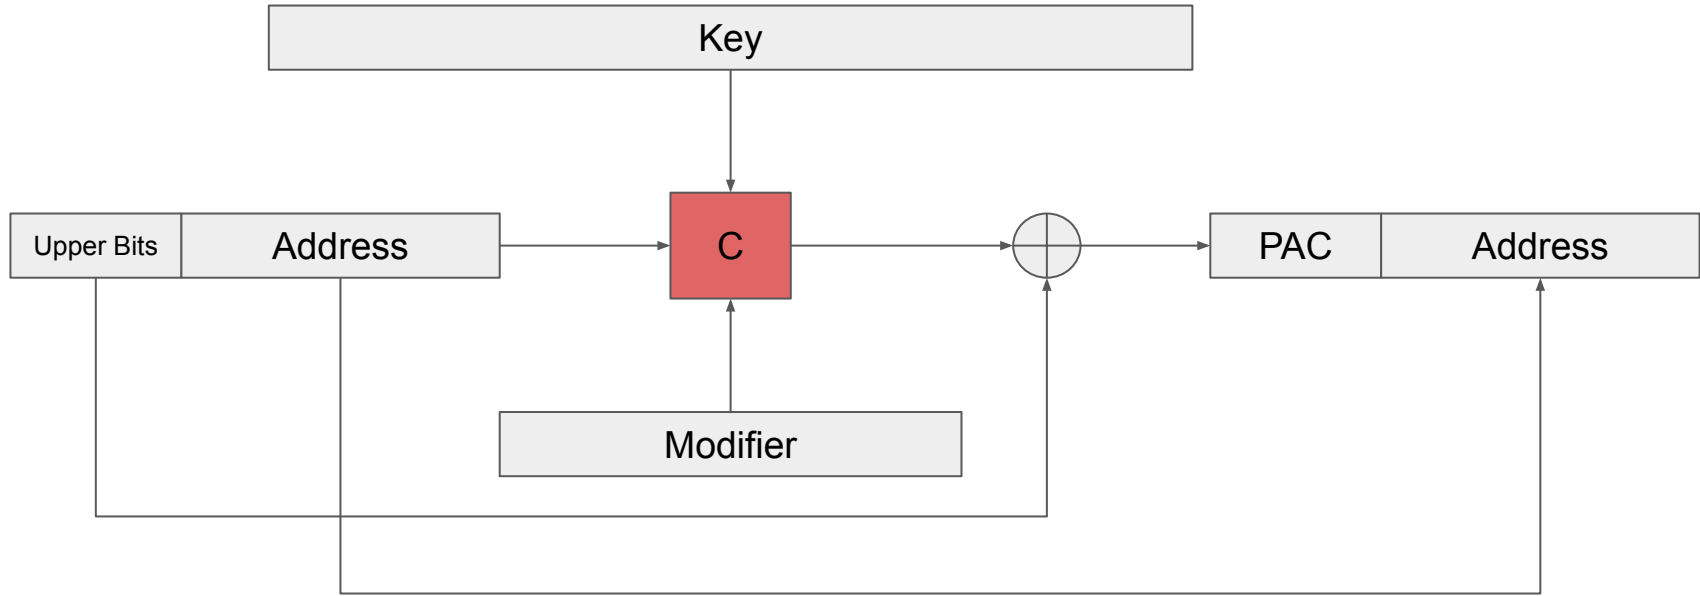
- Use modifier to **restrict** a signed pointer to a specific **context** 💡

$$C(P1, K1, M1) = PAC1$$

$$C(P1, K1, M2) = PAC2$$

- Modifier not secret
- Possible values: stack pointer, register, zero

PAuth: Cipher



PAuth: The QARMA Cipher

Reference cipher developed at Qualcomm's Product Security office in Munich and was named after **Q**ualcomm **ARM** Authenticator.

Cipher:

- Lightweight Tweakable Block Cipher (TBC)
- Purpose-designed with low latency, power consumption, silicon real-estate in mind
- 64-bit and 128-bit block-size / 128-bit and 256-bit key-size

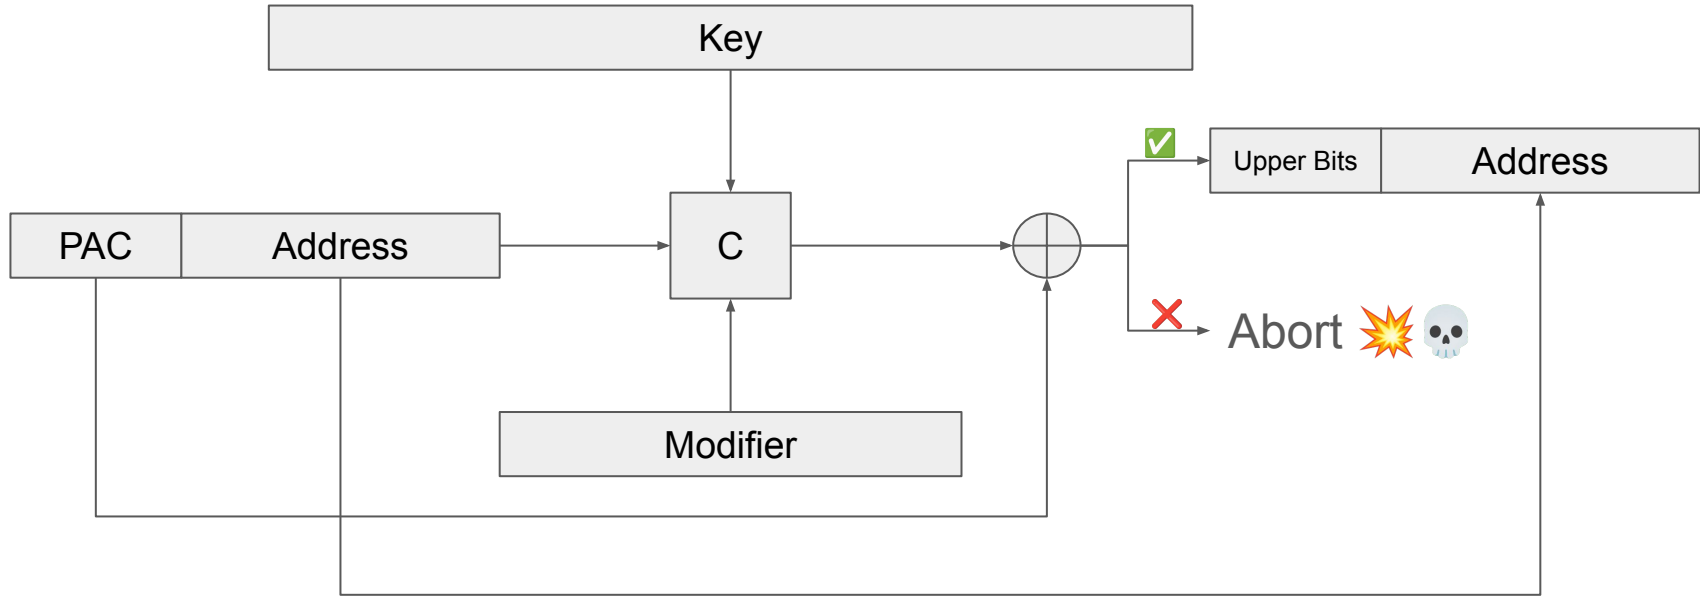
In PAuth:

- 64-bit blocks and 128-bit keys
- Available variants are QARMA-5 and QARMA-3

For more info see:

- **Paper:** <https://eprint.iacr.org/2016/444.pdf>
- **Submission to NIST:**
<https://csrc.nist.gov/CSRC/media/Projects/Lightweight-Cryptography/documents/round-1/spec-doc/gameleon-spec.pdf>

PAuth Operation: Authenticate



PAAuth: Implementation

New Instructions: Basic Set (32 instr.)

Instruction Class	Purpose
PAC [key][modifier]	Sign pointer with key and modifier
AUT [key][modifier]	Authenticate pointer with key and modifier
XPAC [I,D,LRI]	Strip PAC from pointer into register



Implemented as **HINT** instructions, execute as **NOP** on pre-Armv8.3 CPUs.

PAuth: Implementation

New Instructions: Combined set (14 instr)

Instruction Class	Purpose	Key Type
RETA [A,B]	Authenticated return	Instruction Key A, B
BRA [A,B][Z]	Authenticated branch	Instruction Key A, B
BLRA [A,B]Z	Authenticated branch-link	Instruction Key A, B
ERETA [A,B]	Authenticated exception return	Instruction Key A, B
LDRA [A,B]	Authenticated load	Data Key A, B



Implemented as **FUSED** instructions, cause **FAULT** on pre-Armv8.3 CPUs.

PAuth: Implementation

0000000000400564 <main>:

```
400564: d503233f paciasp
400568: a9bf7bfd stp    x29, x30, [sp, #-16]!
40056c: 910003fd mov    x29, sp
400570: 90000000 adrp   x0, 400000 <_init-0x3e8>
400574: 9118c000 add    x0, x0, #0x630
400578: 97ffffb6 bl     400450 <puts@plt>
40057c: 52800000 mov    w0, #0x0
400580: a8c17bfd ldp    x29, x30, [sp], #16
400584: d50323bf autiasp
400588: d65f03c0 ret
```

0000000000400564 <main>:

```
400564: d503233f paciasp
400568: a9bf7bfd stp    x29, x30, [sp, #-16]!
40056c: 910003fd mov    x29, sp
400570: 90000000 adrp   x0, 400000 <_init-0x3e8>
400574: 9118a000 add    x0, x0, #0x628
400578: 97ffffb6 bl     400450 <puts@plt>
40057c: 52800000 mov    w0, #0x0
400580: a8c17bfd ldp    x29, x30, [sp], #16
400584: d65f0bff retaa
```

Branch Target Identification (BTI)

Branch Target Identification (BTI)

Armv8.5-A introduces **FEAT_BTI** as a **MANDATORY** extension.

Mitigates attacks on *indirect jumps* (BLR, BR, RET).

Branch Target Identification (BTI)

```
4010d1e4: eb00007f      cmp     x3, x0
4010d1e8: 54ffffc1      b.ne    4010d1e0 <strerror_r+0xd0>  // b.any
4010d1ec: 17ffffdf      b       4010d168 <strerror_r+0x58>
4010d1f0: d503249f      bti     j
4010d1f4: f0000081      adrp    x1, 40120000 <days_per_mon+0x15c8>
4010d1f8: 912e8021      add     x1, x1, #0xba0
4010d1fc: 17ffffeb      b       4010d1a8 <strerror_r+0x98>
4010d200: d503249f      bti     j
4010d204: f0000081      adrp    x1, 40120000 <days_per_mon+0x15c8>
4010d208: 912de021      add     x1, x1, #0xb78
4010d20c: 17ffffe7      b       4010d1a8 <strerror_r+0x98>
4010d210: d503249f      bti     j
4010d214: f0000081      adrp    x1, 40120000 <days_per_mon+0x15c8>
4010d218: 912d8021      add     x1, x1, #0xb60
4010d21c: 17ffffe3      b       4010d1a8 <strerror_r+0x98>
4010d220: d503249f      bti     j
4010d224: f0000081      adrp    x1, 40120000 <days_per_mon+0x15c8>
4010d228: 912d4021      add     x1, x1, #0xb50
4010d22c: 17ffffdf      b       4010d1a8 <strerror_r+0x98>
4010d230: d503249f      bti     j
...
```


BTI: Implementation

New Instruction: `BTI {type}`

The compiler generates BTI instructions with {target} set to the expected branch type.

- **c**: indirect calls (BLR Xn)
- **j**: indirect jumps (BR Xn)
- **jc**: indirect jumps or indirect calls

 The BTI instruction is implemented on the **HINT** space, will execute as **NOP** on pre-Armv8.5-A CPUs.

BTI: Implementation

Page Tables: New Guarded Page (GP) attribute

- Pages marked as *guarded* are protected with BTI
- Pages marked as *unguarded* don't require a BTI instruction on branch targets

Allows backwards compatibility with code that has not been compiled with BTI.

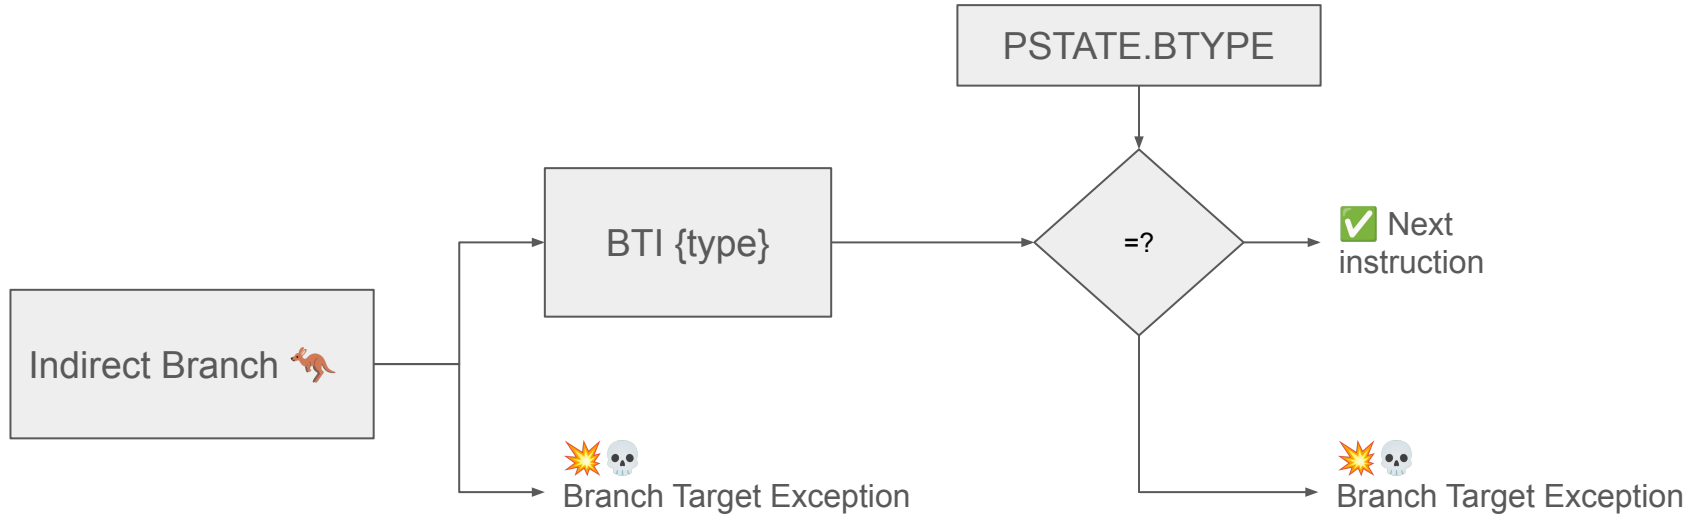
BTI: Implementation

New Process State (PSTATE) Bits: BTYPE

Encodes branch type. Updated on every instruction.

BTYPE	Instructions	Memory Type	Register	Notes
0b00	Non-branch	Any	Any	-
0b01	Any	Non-Guarded	Any	-
0b01	BR, BRAA, BRAAZ, BRAB, BRABZ	Guarded	X16, x17	IP0, IP1 (see APCS) Non-function calls. Veneers, Jump Tables.
0b10	BLR, BLRAA, BLRAAZ, BLRAB, BLRABZ	Any	Any	Function calls.
0b11	BR, BRAA, BRAAZ, BRAB, BRABZ	Guarded	Any except x16, x17	Non-function calls. Case statements.

BTI: Implementation



Ecosystem Support

Compiler Support (gcc / clang)

Common parameter for gcc & clang:

`-mbranch-protection=<type>`

<type> values:

- **none:** Disables all branch protections
- **pac-ret:** Enables PAuth for function returns on non-leaf functions
 - **+leaf** addendum to sign return address of leaf functions
 - **+bkey** addendum to use Key B
- **bti:** Enables BTI
- **standard:** Enables all protections

Compiler Support (gcc / clang)

How good is it?

- PAuth only protects the return address (RoP)
- No C++ vtables protection
- No authenticated branching
- No authenticated loads
- Relies on BTI to protect against JoP

Apple's fork of LLVM is more interesting: <https://github.com/apple/llvm-project>

Linux / Android Support

Userland: CONFIG_ARM64_PTR_AUTH / CONFIG_ARM64_BTI

- Per-process keys (IA, IB ,DA, DB, G). Kernel updates on context-switch.
- Process keys shared between threads.
- Inherit keys upon fork(), re-init keys upon execve().

Kernel: CONFIG_ARM64_PTR_AUTH_KERNEL / CONFIG_ARM64_BTI_KERNEL

- Only IA
- Key switched upon kernel entry / exit.

Again, Apple more interesting.

Hardware Support

SoC	PAuth	BTI
Amazon Graviton 3	Y	N
Amazon Graviton 4	Y	Y
Apple M1	Y	N
Apple M2	Y	Y
Apple M3	Y	Y
Google Tensor G3	Y	Y
Mediatek Dimensity 9000	Y	Y
Qualcomm Snapdragon 8 Gen 1 (SM8450)	Y	Y

SoC	PAuth	BTI
Qualcomm Snapdragon 8 Gen 2 (SM8550)	Y	Y
Qualcomm Snapdragon 8 Gen 3 (SM8650)	Y	Y
Qualcomm Snapdragon 8cx Gen3	Y	N
Samsung Exynos 2200	Y	Y
Samsung Exynos 2400	Y	Y

Evaluation

Evaluation: PAuth Security

Probability to guess PAC:

$$P(\text{💀}) = 1 / 2^{\text{PAC_bits}}$$

Virtual Address Size	TBI disabled MTE disabled	TBI disabled MTE enabled	TBI enabled MTE don't care
36-bits	PAC bits: 27 $P(\text{💀}) = 1 / 134,217,728$	PAC bits: 23 $P(\text{💀}) = 1 / 8,388,608$	PAC bits: 19 $P(\text{💀}) = 1 / 524,288$
39-bits	PAC bits: 24 $P(\text{💀}) = 1 / 16,777,216$	PAC bits: 20 $P(\text{💀}) = 1 / 1,048,576$	PAC bits: 16 $P(\text{💀}) = 1 / 65,536$
42-bits	PAC bits: 21 $P(\text{💀}) = 1 / 2,097,152$	PAC bits: 17 $P(\text{💀}) = 1 / 131,072$	PAC bits: 13 $P(\text{💀}) = 1 / 8,192$
48-bits	PAC bits: 15 $P(\text{💀}) = 1 / 32,768$	PAC bits: 11 $P(\text{💀}) = 1 / 2,048$	PAC bits: 7 $P(\text{💀}) = 1 / 128$
52-bits	PAC bits: 11 $P(\text{💀}) = 1 / 2,048$	PAC bits: 7 $P(\text{💀}) = 1 / 128$	PAC bits: 3 $P(\text{💀}) = 1 / 8$

Evaluation: BTI Security

Given a object with **M** BTI instructions, the probability to find a chain of **N** gadgets is:

$$P(\text{💀}) = M^N$$

For **M** \approx 1%, $P(\text{💀})$ approaches zero as **N** grows.

Evaluation: PAuth Effectiveness

Arm's white-paper on PAuth / BTI [2]:

- *ROP gadgets reduced by ~77x when PAuth is used without BTI.*

Based on glibc (unknown version) & ROPGadget (unknown gadget depth)

Arm's presentation on Chromium [3]:

- *ROP gadgets reduced by ~59x.*
- *Report combined results for ROP/JOP with Pauth+BTI.*

Based on glibc (Ubuntu 14.04) and ROPGadget (gadget depth ≤ 10)

Evaluation: PAuth Effectiveness

My numbers (Nginx 1.15 / Unikraft 0.17.0 / GCC-13)


None:

```
 ROPgadget.py --binary --nojob build/nginx_qemu-arm64
```

 41,386 unique gadgets

PAuth:

 RET → ROP gadget ⇒ RETAA → ! ROP gadget

```
 ROPgadget.py --binary --nojob build/nginx_qemu-arm64 | grep -v retaa
```

 3,285 unique gadgets

Result: 92% improvement

Evaluation: PAuth Effectiveness

How about the remaining 8%?

```
00000000401fb4d0 <timespec_get>:
    401fb4d0: 7100043f      cmp    w1, #0x1
    401fb4d4: 54000060      b.eq   401fb4e0 <timespec_get+0x10> // b.none
    401fb4d8: 52800000      mov    w0, #0x0 // #0
    401fb4dc: d65f03c0      ret
    401fb4e0: d503233f      paciasp
    401fb4e4: a9bf7bfd      stp    x29, x30, [sp, #-16]!
    401fb4e8: aa0003e1      mov    x1, x0
    401fb4ec: 910003fd      mov    x29, sp
    401fb4f0: 52800000      mov    w0, #0x0 // #0
    401fb4f4: 97fff7f3      bl     401f94c0 <__clock_gettime>
    401fb4f8: 2a2003e0      mvn    w0, w0
    401fb4fc: a8c17bfd      ldp    x29, x30, [sp], #16
    401fb500: 531f7c00      lsr    w0, w0, #31
    401fb504: d65f0bff      retaa
```

Evaluation: BTI Effectiveness

My numbers (Nginx 1.15 / Unikraft 0.17.0 / GCC-13 / ROPGadget)

None:



JOP Gadget = #depth instructions to next indirect branch



ROPGadget uses BR, BLR, RET for JOP gadgets \Rightarrow patch to exclude RET



ROPgadget.py --binary --nojop build/nginx_qemu-arm64



7,944 unique gadgets

Evaluation: BTI Effectiveness

```
diff --git a/ropgadget/gadgets.py b/ropgadget/gadgets.py
index a1346bb..a5a4c61 100644
--- a/ropgadget/gadgets.py
+++ b/ropgadget/gadgets.py
@@ -264,8 +264,9 @@ class Gadgets(object):
     ]
     else:
         gadgets = [
-            [b"\x00\x20\x40\x60\x80\xa0\xc0\xe0]{1}[\x00-\x03]{1}[\x1f\x5f]{1}\xd6", 4, 4], # br reg
+            [b"\x00\x20\x40\x60\x80\xa0\xc0\xe0]{1}[\x00-\x03]{1}\x1f\xd6", 4, 4], # br reg excl.ret
+            [b"\x00\x20\x40\x60\x80\xa0\xc0\xe0]{1}[\x00-\x03]{1}\xd6", 4, 4] # blr reg # noqa: W605 # FIXME: \?
+
         ]
         arch_mode = CS_MODE_ARM
     elif arch == CS_ARCH_ARM:
```

Evaluation: BTI Effectiveness

BTI:

💡 JOP Gadget = BTI + #depth instructions to next indirect branch

⚠️ ROPGadget doesn't know about BTI ⇒ Filter out gadgets with BTI

```
ROPgadget.py --norop --binary build/nginx_qemu-arm64 | \
  sed 's/.* : //' | grep '^bti' | sort | uniq | wc -l
```

➡️ 48† unique gadgets

Result: 99.4% improvement

† In practice even less as don't take into account the branch type

Evaluation: PAuth Performance

Instruction Overhead:

+2 instructions per function (PACIASP/RETAA)

Instruction Latency: (Neoverse N2)

- **PACIASP**: 5 cycles
- **RETAA**: 5 cycles

➡ +10 cycles per function

Evaluation: BTI Performance

Instruction Overhead:

+1 instruction per indirect branch (BTI)

Instruction Latency: (Neoverse N2)


- **BTI:** undocumented (should be 1 cycle)

➡ +1 cycle per indirect branch


Evaluation: Performance

My numbers:

nbench-byte



wrk nginx




Unikraft 0.17.0 / GCC-13 / QEMU 9.x with HVF acceleration on Apple M2 (Mac Mini)

Evaluation: Code Size

Arm's measurements [2]: glibc (unknown version)

- PAuth (basic set): +1.77%
- BTI (basic set): +1.22%
- PAuth + BTI (basic set): +2.90%
- Pauth + BTI (comb. set): +1.61%




Evaluation: Code Size

My numbers:

nginx 1.25 / Unikraft 0.17.0 / GCC-13











- **PAuth (comb. set): +1.52%**
- **BTI (comb. set): +1.26%**
- **Pauth + BTI (comb. set): +1.77%**



	None	PAuth	BTI	PAuth + BTI
text	1311264	1333316	1331652	1337188
total	2094704	2123376	2115184	2127472
increase %	0%	1.52%	1.26%	1.77%

In Practice

Surely CFI is enforced by now!  

-  [Examining Pointer Authentication on the iPhone XS](#)
-  [Splitting atoms in XNU](#)
-  [HackPac: Hacking Pointer Authentication in iOS User Space](#)
-  [iOS Kernel PAC, One Year Later](#)
-  [Fugu 14 jailbreak](#)
-  [Everything has Changed in iOS 14, but Jailbreak is Eternal](#)
-  [Fugu 15 jailbreak](#)
-  [PACMAN: Attacking ARM Pointer Authentication with Speculative Execution](#)
-  [Apple PAC, Four Years Later](#)
-  [Demystifying Pointer Authentication on Apple M1](#)

Final Thoughts

- PAuth & BTI are performant and effective.
- Great value for money: Implementation effort vs improvement.
- Pragmatic: No code rewrite in applications, minimal implementation in kernel.
- A leap forward compared to software mitigations.
- Eventually expected to be available in all Arm-based SoC.
- Although not impossible to bypass, still make CFI attacks significantly harder.
- More hardware-assisted mitigations on the way:
 - o Memory Tagging (MTE) against memory safety violations.
 - o Keep an eye on CHERI.

QUESTIONS